Bibliography

Abramowitz, M., and I. A. Stegun. 1965. *Handbook of Mathematical Functions.* Dover.

Adler, S. L. 1981. "Over-relaxation method for the Monte Carlo evaluation of the partition function for multiquadratic actions." *Physical Review D* 23:2901–2904.

Aghajanyan, Armen, Bernie Huang, Candace Ross, Vladimir Karpukhin, Hu Xu, Naman Goyal, Dmytro Okhonko, *et al.* 2022. *CM3: A Causal Masked Multimodal Model of the Internet.* Technical report. arXiv:2201.07520.

Aghajanyan, Armen, Luke Zettlemoyer, and Sonal Gupta. 2020. *Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning.* Technical report. arXiv:2012.13255.

Ahn, J. H., and J. H. Oh. 2003. "A constrained EM algorithm for principal component analysis." *Neural Computation* 15 (1): 57–65.

Alayrac, Jean-Baptiste, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, *et al.* 2022. *Flamingo: a Visual Language Model for Few-Shot Learning.* Technical report. arXiv:2204.14198.

Amari, S., A. Cichocki, and H. H. Yang. 1996. "A new learning algorithm for blind signal separation." In *Advances in Neural Information Processing Systems,* edited by D. S. Touretzky, M. C. Mozer, and M. E. Hasselmo, 8:757–763. MIT Press.

Anderson, J. A., and E. Rosenfeld. 1988. *Neurocomputing: Foundations of Research.* MIT Press.

Anderson, T. W. 1963. "Asymptotic Theory for Principal Component Analysis." *Annals of Mathematical Statistics* 34:122–148.

Arjovsky, M., S. Chintala, and L. Bottou. 2017. *Wasserstein GAN.* Technical report. arXiv:1701.07875.

Attias, H. 1999. "Independent factor analysis." *Neural Computation* 11 (4): 803–851.

Austin, Jacob, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. 2021. "Structured Denoising Diffusion Models in Discrete State-Spaces." In *Advances in Neural Information Processing Systems,* 34:17981–17993.

Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E Hinton. 2016. *Layer Normalization.* Technical report. arXiv:1607.06450.

Bach, F. R., and M. I. Jordan. 2002. "Kernel Independent Component Analysis." *Journal of Machine Learning Research* 3:1–48.

Badrinarayanan, Vijay, Alex Kendall, and Roberto Cipolla. 2015. *SegNet: A Deep Convolu-*

© The Author(s), under exclusive license to Springer Nature Switzerland AG 2024
C. M. Bishop, H. Bishop, *Deep Learning*, https://doi.org/10.1007/978-3-031-45468-4

tional Encoder-Decoder Architecture for Image Segmentation. Technical report. arXiv:1511.00561.

Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. 2014. *Neural Machine Translation by Jointly Learning to Align and Translate.* Technical report. arXiv:1409.0473.

Baldi, P., and K. Hornik. 1989. "Neural networks and principal component analysis: learning from examples without local minima." *Neural Networks* 2 (1): 53–58.

Balduzzi, David, Marcus Frean, Lennox Leary, JP Lewis, Kurt Wan-Duo Ma, and Brian McWilliams. 2017. *The Shattered Gradients Problem: If resnets are the answer, then what is the question?* Technical report. arXiv:1702.08591.

Bartholomew, D J. 1987. *Latent Variable Models and Factor Analysis.* Charles Griffin.

Basilevsky, Alexander. 1994. *Statistical Factor Analysis and Related Methods: Theory and Applications.* Wiley.

Bather, J. 2000. *Decision Theory: An Introduction to Dynamic Programming and Sequential Decisions.* Wiley.

Battaglia, Peter W., Jessica B. Hamrick, Victor Bapst, Alvaro Sanchez-Gonzalez, Vinicius Zambaldi, Mateusz Malinowski, Andrea Tacchetti, *et al.* 2018. *Relational inductive biases, deep learning, and graph networks.* Technical report. arXiv:1806.01261.

Baydin, A. G., B. A. Pearlmutter, A. A. Radul, and J. M. Siskind. 2018. "Automatic differentiation in machine learning: a survey." *Journal of Machine Learning Research* 18:1–43.

Becker, S., and Y. LeCun. 1989. "Improving the convergence of back-propagation learning with second order methods." In *Proceedings of the 1988 Connectionist Models Summer School,* edited by D. Touretzky, G. E. Hinton, and T. J. Sejnowski, 29–37. Morgan Kaufmann.

Belkin, Mikhail, Daniel Hsu, Siyuan Ma, and Soumik Mandal. 2019. "Reconciling modern machine-learning practice and the classical bias-variance trade-off." *Proceedings of the National Academy of Sciences* 116 (32): 15849–15854.

Bell, A. J., and T. J. Sejnowski. 1995. "An information maximization approach to blind separation and blind deconvolution." *Neural Computation* 7 (6): 1129–1159.

Bellman, R. 1961. *Adaptive Control Processes: A Guided Tour.* Princeton University Press.

Bengio, Yoshua, Aaron Courville, and Pascal Vincent. 2012. *Representation Learning: A Review and New Perspectives.* Technical report. arXiv:1206.5538.

Bengio, Yoshua, Nicholas Léonard, and Aaron Courville. 2013. *Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation.* Technical report. arXiv:1308.3432.

Berger, J. O. 1985. *Statistical Decision Theory and Bayesian Analysis.* Second. Springer.

Bernardo, J. M., and A. F. M. Smith. 1994. *Bayesian Theory.* Wiley.

Bishop, C. M. 1995a. "Regularization and Complexity Control in Feed-forward Networks." In *Proceedings International Conference on Artificial Neural Networks ICANN'95,* edited by F. Fougelman-Soulie and P. Gallinari, 1:141–148. EC2 et Cie.

Bishop, Christopher M. 1992. "Exact Calculation of the Hessian Matrix for the Multilayer Perceptron." *Neural Computation* 4 (4): 494–501.

Bishop, Christopher M. 1994. "Novelty Detection and Neural Network Validation." *IEE Proceedings: Vision, Image and Signal Processing* 141 (4): 217–222.

Bishop, Christopher M. 1995b. *Neural Networks for Pattern Recognition.* Oxford University Press.

Bishop, Christopher M. 1995c. "Training with noise is equivalent to Tikhonov regularization." *Neural Computation* 7 (1): 108–116.

Bishop, Christopher M. 2006. *Pattern Recognition and Machine Learning.* Springer.

Bommasani, Rishi, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, *et al.* 2021. *On the Opportunities and Risks of Foundation Models.* Technical report. arXiv:2108.07258.

Bottou, L. 2010. "Large-scale machine learning with stochastic gradient descent." In *Proceedings COMPSTAT 2010,* 177–186. Springer.

Bourlard, H., and Y. Kamp. 1988. "Auto-association by multilayer perceptrons and singular value decomposition." *Biological Cybernetics* 59:291–294.

Breiman, L. 1996. "Bagging predictors." *Machine Learning* 26:123–140.

Brinker, T. J., A. Hekler, A. H. Enk, C. Berking, S Haferkamp, A. Hauschild, M. Weichenthal, *et al.* 2019. "Deep neural networks are superior to dermatologists in melanoma image classification." *European Journal of Cancer* 119:11–17.

Brock, Andrew, Jeff Donahue, and Karen Simonyan. 2018. "Large-Scale GAN Training for High Fidelity Natural Image Synthesis." In *Proceedings of the International Conference Learning Representations (ICLR).* ArXiv:1809.11096.

Bronstein, Michael M., Joan Bruna, Taco Cohen, and Petar Velickovic. 2021. *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges.* Technical report. arXiv:2104.13478.

Bronstein, Michael M., Joan Bruna, Yann LeCun, Arthur Szlam, and Pierre Vandergheynst. 2017. "Geometric Deep Learning: Going Beyond Eulcidean Data." In *IEEE Signal Processing Magazine,* vol. 34. 4. IEEE, July.

Broomhead, D. S., and D. Lowe. 1988. "Multivariable functional interpolation and adaptive networks." *Complex Systems* 2:321–355.

Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, *et al.* 2020. *Language Models are Few-Shot Learners.* Technical report. arXiv:2005.14165.

Bubeck, Sébastien, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, *et al.* 2023. *Sparks of Artificial General Intelligence: Early experiments with GPT-4.* Technical report. arXiv:2303.12712.

Cardoso, J-F. 1998. "Blind signal separation: statistical principles." *Proceedings of the IEEE* 9 (10): 2009–2025.

Caruana, R. 1997. "Multitask learning." *Machine Learning* 28:41–75.

Casella, G., and R. L. Berger. 2002. *Statistical Inference.* Second. Duxbury.

Chan, K., T. Lee, and T. J. Sejnowski. 2003. "Variational Bayesian learning of ICA with missing data." *Neural Computation* 15 (8): 1991–2011.

Chen, A. M., H. Lu, and R. Hecht-Nielsen. 1993. "On the geometry of feedforward neural network error surfaces." *Neural Computation* 5 (6): 910–927.

Chen, Mark, Alec Radford, Rewon Child, Jeffrey Wu, Heewoo Jun, David Luan, and Ilya Sutskever. 2020. "Generative Pretraining From Pixels." *Proceedings of Machine Learning Research* 119:1691–1703.

Chen, R. T. Q., Rubanova Y, J. Bettencourt, and D. Duvenaud. 2018. *Neural Ordinary Differential Equations.* Technical report. arXiv:1806.07366.

Chen, Ting, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. *A Simple Framework for Contrastive Learning of Visual Representations.* Technical report. arXiv:2002.05709.

Cho, Kyunghyun, Bart van Merrienboer, Çaglar Gülçehre, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. 2014. *Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translations.* Technical report. arXiv:1406.1078.

Choudrey, R. A., and S. J. Roberts. 2003. "Variational mixture of Bayesian independent component analyzers." *Neural Computation* 15 (1): 213–252.

Christiano, Paul, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. *Deep reinforcement learning from human preferences.* Technical report. arXiv:1706.03741.

Collobert, R. 2004. "Large Scale Machine Learning." PhD diss., Université Paris VI.

Comon, P., C. Jutten, and J. Herault. 1991. "Blind source separation, 2: problems statement." *Signal Processing* 24 (1): 11–20.

Cover, T., and P. Hart. 1967. "Nearest neighbor pattern classification." *IEEE Transactions on Information Theory* IT-11:21–27.

Cover, T. M., and J. A. Thomas. 1991. *Elements of Information Theory.* Wiley.

Cox, R. T. 1946. "Probability, frequency and reasonable expectation." *American Journal of Physics* 14 (1): 1–13.

Cybenko, G. 1989. "Approximation by superpositions of a sigmoidal function." *Mathematics of Control, Signals and Systems* 2:304–314.

Dawid, A. P. 1979. "Conditional Independence in Statistical Theory (with discussion)." *Journal of the Royal Statistical Society, Series B* 4:1–31.

Dawid, A. P. 1980. "Conditional Independence for Statistical Operations." *Annals of Statistics* 8:598–617.

Deisenroth, M. P., A. A. Faisal, and C. S. Ong. 2020. *Mathematics for Machine Learning.* Cambridge University Press.

Dempster, A. P., N. M. Laird, and D. B. Rubin. 1977. "Maximum likelihood from incomplete data via the EM algorithm." *Journal of the Royal Statistical Society, B* 39 (1): 1–38.

Deng, Jia, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. "ImageNet: A large-scale hierarchical image database." In *IEEE Conference on Computer Vision and Pattern Recognition.*

Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.* Technical report. arXiv:1810.04805.

Dhariwal, Prafulla, and Alex Nichol. 2021. *Diffusion Models Beat GANs on Image Synthesis.* Technical report. arXiv:2105.05233.

Dinh, Laurent, David Krueger, and Yoshua Bengio. 2014. *NICE: Non-linear Independent Components Estimation.* Technical report. arXiv:1410.8516.

Dinh, Laurent, Jascha Sohl-Dickstein, and Samy Bengio. 2016. *Density estimation using Real NVP.* Technical report. arXiv:1605.08803.

Dodge, Samuel, and Lina Karam. 2017. *A Study and Comparison of Human and Deep Learning Recognition Performance Under Visual Distortions.* Technical report. arXiv:1705.02498.

Doersch, C. 2016. *Tutorial on Variational Autoencoders.* Technical report. arXiv:1606.05908.

Dosovitskiy, Alexey, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, *et al.* 2020. *An Image is Worth* 16×16 *Words: Transformers for Image Recognition at Scale.* Technical report. arXiv:2010.11929.

Duane, S., A. D. Kennedy, B. J. Pendleton, and D. Roweth. 1987. "Hybrid Monte Carlo." *Physics Letters B* 195 (2): 216–222.

Duchi, J., E. Hazan, and Y. Singer. 2011. "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization." *Journal of Machine Learning Research* 12:2121–2159.

Duda, R. O., and P. E. Hart. 1973. *Pattern Classification and Scene Analysis.* Wiley.

Dufter, Philipp, Martin Schmitt, and Hinrich Schütze. 2021. *Position Information in Transformers: An Overview.* Technical report. arXiv:2102.11090.

Dumoulin, Vincent, and Francesco Visin. 2016. *A guide to convolution arithmetic for deep learning.* Technical report. arXiv:1603.07285.

Elliott, R. J., L. Aggoun, and J. B. Moore. 1995. *Hidden Markov Models: Estimation and Control.* Springer.

Esser, Patrick, Robin Rombach, and Björn Ommer. 2020. *Taming Transformers for High-Resolution Image Synthesis.* Technical report. arXiv:2012.09841.

Esteva, A., B. Kuprel, R. A. Novoa, J. Ko, S. M. Swetter, H. M. Blau, and S. Thrun. 2017. "Dermatologist-level classification of skin cancer with deep neural networks." *Nature* 542:115–118.

Everitt, B. S. 1984. *An Introduction to Latent Variable Models.* Chapman / Hall.

Eykholt, Kevin, Ivan Evtimov, Earlence Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul Prakash, Tadayoshi Kohno, and Dawn Song. 2018. "Robust Physical-World Attacks on Deep Learning Visual Classification." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).*

Fawcett, T. 2006. "An introduction to ROC analysis." *Pattern Recognition Letters* 27:861–874.

Feller, W. 1966. *An Introduction to Probability Theory and its Applications.* Second. Vol. 2. Wiley.

Fletcher, R. 1987. *Practical Methods of Optimization.* Second. Wiley.

Forsyth, D. A., and J. Ponce. 2003. *Computer Vision: A Modern Approach.* Prentice Hall.

Freund, Y., and R. E. Schapire. 1996. "Experiments with a new boosting algorithm." In *Thirteenth International Conference on Machine Learning,* edited by L. Saitta, 148–156. Morgan Kaufmann.

Fukushima, K. 1980. "Neocognitron: A Self-organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position." *Biological Cybernetics* 36:193–202.

Funahashi, K. 1989. "On the approximate realization of continuous mappings by neural networks." *Neural Networks* 2 (3): 183–192.

Fung, R., and K. C. Chang. 1990. "Weighting and Integrating Evidence for Stochastic Simulation in Bayesian Networks." In *Uncertainty in Artificial Intelligence,* edited by P. P. Bonissone, M. Henrion, L. N. Kanal, and J. F. Lemmer, 5:208–219. Elsevier.

Gatys, Leon A., Alexander S. Ecker, and Matthias Bethge. 2015. *A Neural Algorithm of Artistic Style.* Technical report. arXiv:1508.06576.

Geman, S., and D. Geman. 1984. "Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images." *IEEE PAMI* 6 (1): 721–741.

Gemmeke, Jort F., Daniel P. W. Ellis, Dylan Freedman, Aren Jansen, Wade Lawrence, R. Channing Moore, Manoj Plakal, and Marvin Ritter. 2017. "Audio Set: An ontology and human-labeled dataset for audio events." In *Proc. IEEE ICASSP 2017.* New Orleans, LA.

Germain, Mathieu, Karol Gregor, Iain Murray, and Hugo Larochelle. 2015. *MADE: Masked Autoencoder for Distribution Estimation.* Technical report. arXiv:1502.03509.

Gilks, W. R. 1992. "Derivative-free adaptive rejection sampling for Gibbs sampling." In *Bayesian Statistics,* edited by J. Bernardo, J. Berger, A. P. Dawid, and A. F. M. Smith, vol. 4. Oxford University Press.

Gilks, W. R., N. G. Best, and K. K. C. Tan. 1995. "Adaptive rejection Metropolis sampling." *Applied Statistics* 44:455–472.

Gilks, W. R., S. Richardson, and D. J. Spiegelhalter. 1996. *Markov Chain Monte Carlo in Practice.* Chapman / Hall.

Gilks, W. R., and P. Wild. 1992. "Adaptive rejection sampling for Gibbs sampling." *Applied Statistics* 41:337–348.

Gilmer, Justin, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. 2017. *Neural Message Passing for Quantum Chemistry.* Technical report. arXiv:1704.01212.

Girshick, Ross B. 2015. *Fast R-CNN.* Technical report. arXiv:1504.08083.

Golub, G. H., and C. F. Van Loan. 1996. *Matrix Computations.* Third. John Hopkins University Press.

Gong, Yuan, Yu-An Chung, and James R. Glass. 2021. *AST: Audio Spectrogram Transformer.* Technical report. arXiv:2104.01778.

Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning.* MIT Press.

Goodfellow, Ian J., Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. *Generative Adversarial Networks.* Technical report. arXiv:1406.2661.

Goodfellow, Ian J., Jonathon Shlens, and Christian Szegedy. 2014. *Explaining and Harnessing Adversarial Examples.* Technical report. arXiv:1412.6572.

Grathwohl, Will, Ricky T. Q. Chen, Jesse Bettencourt, Ilya Sutskever, and David Duvenaud. 2018. *FFJORD: Free-form Continuous Dynamics for Scalable Reversible Generative Models.* Technical report. arXiv:1810.01367.

Griewank, A., and A Walther. 2008. *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation.* Second. SIAM.

Grosse, R. 2018. *Automatic Differentiation.* CSC321 Lecture 10. University of Toronto.

Gulrajani, I., F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville. 2017. *Improved training of Wasserstein GANs.* Technical report. arXiv:1704.00028.

Gutmann, Michael, and Aapo Hyvärinen. 2010. "Noise-contrastive estimation: A new estimation principle for unnormalized statistical models." *Journal of Machine Learning Research* 9:297–304.

Hamilton, W. L. 2020. *Graph Representation Learning.* Morgan / Claypool.

Hartley, R., and A. Zisserman. 2004. *Multiple View Geometry in Computer Vision.* Second. Cambridge University Press.

Hassibi, B., and D. G. Stork. 1993. "Second order derivatives for network pruning: optimal brain surgeon." In *Proceedings International Conference on Neural Information Processing Systems (NeurIPS),* edited by S. J. Hanson, J. D. Cowan, and C. L. Giles, 5:164–171. Morgan Kaufmann.

Hastie, T., R. Tibshirani, and J. Friedman. 2009. *The Elements of Statistical Learning.* Second. Springer.

Hastings, W. K. 1970. "Monte Carlo sampling methods using Markov chains and their applications." *Biometrika* 57:97–109.

He, Kaiming, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross B. Girshick. 2021. *Masked Autoencoders Are Scalable Vision Learners.* Technical report. arXiv:2111.06377.

He, Kaiming, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. 2019. *Momentum Contrast for Unsupervised Visual Representation Learning.* Technical report. arXiv:1911.05722.

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2015a. *Deep Residual Learning for Image Recognition.* Technical report. arXiv:1512.03385.

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2015b. *Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification.* Technical report. arXiv:1502.01852.

Henrion, M. 1988. "Propagation of Uncertainty by Logic Sampling in Bayes' Networks." In *Uncertainty in Artificial Intelligence,* edited by J. F. Lemmer and L. N. Kanal, 2:149–164. North Holland.

Higgins, I., L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinik, S. Mohamed, and A. Lerchner. 2017. "β-VAE: learning basic visual concepts with a constrained variational framework." In *Proceedings of the International Conference Learning Representations (ICLR).*

Hinton, G. E. 2012. *Neural Networks for Machine Learning. Lecture 6.5. Coursera Lectures.*

Hinton, G. E., M. Welling, Y. W. Teh, and S Osindero. 2001. "A new view of ICA." In *Proceedings of the International Conference on Independent Component Analysis and Blind Signal Separation,* vol. 3.

Hinton, Geoffrey, Oriol Vinyals, and Jeff Dean. 2015. *Distilling the Knowledge in a Neural Network.* Technical report. arXiv:1503.02531.

Hinton, Geoffrey E. 2002. "Training products of experts by minimizing contrastive divergence." *Neural Computation* 14:1771–1800.

Ho, Jonathan, Ajay Jain, and Pieter Abbeel. 2020. *Denoising Diffusion Probabilistic Models.* Technical report. arXiv:2006.11239.

Ho, Jonathan, Chitwan Saharia, William Chan, David J. Fleet, Mohammad Norouzi, and Tim Salimans. 2021. *Cascaded Diffusion Models for High Fidelity Image Generation.* Technical report. arXiv:2106.15282.

Hochreiter, S., and J. Schmidhuber. 1997. "Long short-term Memory." *Neural Computation* 9 (8): 1735–1780.

Hojen-Sorensen, P. A., O. Winther, and L. K. Hansen. 2002. "Mean field approaches to independent component analysis." *Neural Computation* 14 (4): 889–918.

Holtzman, Ari, Jan Buys, Maxwell Forbes, and Yejin Choi. 2019. *The Curious Case of Neural Text Degeneration.* Technical report. arXiv:1904.09751.

Hornik, K., M. Stinchcombe, and H. White. 1989. "Multilayer feedforward networks are universal approximators." *Neural Networks* 2 (5): 359–366.

Hospedales, Timothy, Antreas Antoniou, Paul Micaelli, and Amos Storkey. 2021. "Meta-learning in neural networks: A survey." *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44 (9): 5149–5169.

Hotelling, H. 1933. "Analysis of a complex of statistical variables into principal components." *Journal of Educational Psychology* 24:417–441.

Hotelling, H. 1936. "Relations between two sets of variables." *Biometrika* 28:321–377.

Hu, Anthony, Lloyd Russell, Hudson Yeo, Zak Murez, George Fedoseev, Alex Kendall, Jamie Shotton, and Gianluca Corrado. 2023. *GAIA-1: A Generative World Model for Autonomous Driving.* Technical report. arXiv:2309.17080.

Hu, Edward J., Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. *LoRA: Low-Rank Adaptation of Large Language Models.* Technical report. arXiv:2106.09685.

Hubel, D. H., and T. N. Wiesel. 1959. "Receptive fields of single neurons in the cat's striate cortex." *Journal of Physiology* 148:574–591.

Hyvärinen, A. 2005. "Estimation of Non-Normalized Statistical Models by Score Matching." *Journal of Machine Learning Research* 6:695–709.

Hyvärinen, A., and E. Oja. 1997. "A fast fixed-point algorithm for independent component analysis." *Neural Computation* 9 (7): 1483–1492.

Hyvärinen, Aapo, Jarmo Hurri, and Patrick O. Hoyer. 2009. *Natural Image Statistics: A Probabilistic Approach to Early Computational Vision.* Springer.

Ioffe, S., and C. Szegedy. 2015. "Batch normalization." In *Proceedings of the International Conference on Machine Learning (ICML),* 448–456.

Jacobs, R. A., M. I. Jordan, S. J. Nowlan, and G. E. Hinton. 1991. "Adaptive mixtures of local experts." *Neural Computation* 3 (1): 79–87.

Jebara, T. 2004. *Machine Learning: Discriminative and Generative.* Kluwer.

Jensen, C., A. Kong, and U. Kjaerulff. 1995. "Blocking Gibbs sampling in very large probabilistic expert systems." *International Journal of Human Computer Studies. Special Issue on Real-World Applications of Uncertain Reasoning.* 42:647–666.

Jolliffe, I. T. 2002. *Principal Component Analysis.* Second. Springer.

Jumper, John, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, and Olaf Ronneberger. 2021. "Highly accurate protein structure prediction with AlphaFold." *Nature* 596:583–589.

Jutten, C., and J. Herault. 1991. "Blind separation of sources, 1: An adaptive algorithm based on neuromimetic architecture." *Signal Processing* 24 (1): 1–10.

Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. *Scaling Laws for Neural Language Models.* Technical report. arXiv:2001.08361.

Karras, Tero, Timo Aila, Samuli Laine, and Jaakko Lehtinen. 2017. *Progressive Growing of GANs for Improved Quality, Stability, and Variation.* Technical report. arXiv:1710.10196.

Karush, W. 1939. "Minima of functions of several variables with inequalities as side constraints." Master's thesis, Department of Mathematics, University of Chicago.

Khosla, Prannay, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. 2020. *Supervised Contrastive Learning.* Technical report. arXiv:2004.11362.

Kingma, D., and J. Ba. 2014. *Adam: A method for stochastic optimization.* Technical report. arXiv:1412.6980.

Kingma, D. P., and M. Welling. 2013. "Auto-encoding variational Bayes." In *Proceedings of the International Conference on Machine Learning (ICML).* ArXiv:1312.6114.

Kingma, Diederik P., and Max Welling. 2019. *An Introduction to Variational Autoencoders.* Technical report. arXiv:1906.02691.

Kingma, Durk P, Tim Salimans, Rafal Jozefowicz, Xi Chen, Ilya Sutskever, and Max Welling. 2016. "Improved variational inference with inverse autoregressive flow." *Advances in Neural Information Processing Systems* 29.

Kipf, Thomas N., and Max Welling. 2016. *Semi-Supervised Classification with Graph Convolutional Networks.* Technical report. arXiv:1609.02907.

Kloeden, Peter E, and Eckhard Platen. 2013. *Numerical solution of stochastic differential equations.* Vol. 23. Stochastic Modelling and Applied Probability. Springer.

Kobyzev, I., S. J. D. Prince, and M. A. Brubaker. 2019. "Normalizing flows: an introduction and review of current methods." *IEEE Transactions on Pattern Analysis and Machine Intelligence* 43 (11): 3964–3979.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. 2012. "Imagenet classification with deep convolutional neural networks." In *Advances in Neural Information Processing Systems,* vol. 25.

Kuhn, H. W., and A. W. Tucker. 1951. "Nonlinear programming." In *Proceedings of the 2nd Berkeley Symposium on Mathematical Statistics and Probabilities,* 481–492. University of California Press.

Kullback, S., and R. A. Leibler. 1951. "On information and sufficiency." *Annals of Mathematical Statistics* 22 (1): 79–86.

Kurková, V., and P. C. Kainen. 1994. "Functionally Equivalent Feed-forward Neural Networks." *Neural Computation* 6 (3): 543–558.

Lasserre, J., Christopher M. Bishop, and T. Minka. 2006. "Principled hybrids of generative and discriminative models." In *Proceedings 2006 IEEE Conference on Computer Vision and Pattern Recognition, New York.*

Lauritzen, S. L. 1996. *Graphical Models.* Oxford University Press.

Lawley, D. N. 1953. "A Modified Method of Estimation in Factor Analysis and Some Large Sample Results." In *Uppsala Symposium on Psychological Factor Analysis,* 35–42. Number 3 in Nordisk Psykologi Monograph Series. Uppsala: Almqvist / Wiksell.

Lazarsfeld, P. F., and N. W. Henry. 1968. *Latent Structure Analysis.* Houghton Mifflin.

LeCun, Y., B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel. 1989. "Backpropagation Applied to Handwritten ZIP Code Recognition." *Neural Computation* 1 (4): 541–551.

LeCun, Y., L. Bottou, Y. Bengio, and P. Haffner. 1998. "Gradient-Based Learning Applied to Document Recognition." *Proceedings of the IEEE* 86:2278–2324.

LeCun, Y., J. S. Denker, and S. A. Solla. 1990. "Optimal Brain Damage." In *Proceedings International Conference on Neural Information Processing Systems (NeurIPS),* edited by D. S. Touretzky, 2:598–605. Morgan Kaufmann.

LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton. 2015. "Deep Learning." *Nature* 512:436–444.

LeCun, Yann, Sumit Chopra, Raia Hadsell, Marc'Aurelio Ranzato, and Fu-Jie Huang. 2006. "A Tutorial on Energy-Based Learning." In *Predicting Structured Data,* edited by G. Bakir, T. Hofman, B. Schölkopf, A. Smola, and B. Taskar. MIT Press.

Leen, T. K. 1995. "From data distributions to regularization in invariant learning." *Neural Computation* 7:974–981.

Leshno, M., V. Y. Lin, A. Pinkus, and S. Schocken. 1993. "Multilayer feedforward networks with a polynomial activation function can approximate any function." *Neural Networks* 6:861–867.

Li, Hao, Zheng Xu, Gavin Taylor, Christoph Studer, and Tom Goldstein. 2017. *Visualizing the Loss Landscape of Neural Nets.* Technical report. arXiv:1712.09913.

Li, Junnan, Dongxu Li, Caiming Xiong, and Steven Hoi. 2022. *BLIP: Bootstrapping Language-Image Pre-training for Unified Vision-Language Understanding and Generation.* Technical report. arXiv:2201.12086.

Lin, Min, Qiang Chen, and Shuicheng Yan. 2013. *Network in Network.* Technical report. arXiv:1312.4400.

Lin, Tianyang, Yuxin Wang, Xiangyang Liu, and Xipeng Qiu. 2021. *A Survey of Transformers.* Technical report. arXiv:2106.04554.

Lipman, Yaron, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. 2022. *Flow Matching for Generative Modeling.* Technical report arXiv:2210.02747. https://arxiv.org/.

Liu, Pengfei, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2021. *Pre-train, Prompt, and Predict: A Systematic Survey of Prompting Methods in Natural Language Processing.* Technical report. arXiv:2107.13586.

Lloyd, S. P. 1982. "Least squares quantization in PCM." *IEEE Transactions on Information Theory* 28 (2): 129–137.

Long, Jonathan, Evan Shelhamer, and Trevor Darrell. 2014. *Fully Convolutional Networks for Semantic Segmentation.* Technical report. arXiv:1411.4038.

Luo, Calvin. 2022. *Understanding Diffusion Models: A Unified Perspective.* Technical report. arXiv:2208.11970.

Lütkepohl, H. 1996. *Handbook of Matrices.* Wiley.

MacKay, D. J. C. 1992. "A Practical Bayesian Framework for Back-propagation Networks." *Neural Computation* 4 (3): 448–472.

MacKay, D. J. C. 2003. *Information Theory, Inference and Learning Algorithms.* Cambridge University Press.

MacQueen, J. 1967. "Some methods for classification and analysis of multivariate observations." In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability,* edited by L. M. LeCam and J. Neyman, I:281–297. University of California Press.

Magnus, J. R., and H. Neudecker. 1999. *Matrix Differential Calculus with Applications in Statistics and Econometrics.* Wiley.

Mallat, S. 1999. *A Wavelet Tour of Signal Processing.* Second. Academic Press.

Mao, X., Q. Li, H. Xie, R. Lau, Z. Wang, and S. Smolley. 2016. *Least Squares Generative Adversarial Networks.* Technical report. arXiv:1611.04076.

Mardia, K. V., and P. E. Jupp. 2000. *Directional Statistics.* Wiley.

Martens, James, Ilya Sutskever, and Kevin Swersky. 2012. "Estimating the Hessian by Back-propagating Curvature." In *Proceedings of the*

International Conference on Machine Learning (ICML). ArXiv:1206.6464.

McCullagh, P., and J. A. Nelder. 1989. *Generalized Linear Models.* Second. Chapman / Hall.

McCulloch, W. S., and W. Pitts. 1943. "A Logical Calculus of the Ideas Immanent in Nervous Activity." Reprinted in Anderson and Rosenfeld (1988), *Bulletin of Mathematical Biophysics* 5:115–133.

McLachlan, G. J., and T. Krishnan. 1997. *The EM Algorithm and its Extensions.* Wiley.

McLachlan, G. J., and D. Peel. 2000. *Finite Mixture Models.* Wiley.

Meng, X. L., and D. B. Rubin. 1993. "Maximum likelihood estimation via the ECM algorithm: a general framework." *Biometrika* 80:267–278.

Mescheder, L., A. Geiger, and S. Nowozin. 2018. *Which Training Methods for GANs do actually Converge?* Technical report. arXiv:1801.04406.

Metropolis, N., A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller. 1953. "Equation of State Calculations by Fast Computing Machines." *Journal of Chemical Physics* 21 (6): 1087–1092.

Metropolis, N., and S. Ulam. 1949. "The Monte Carlo method." *Journal of the American Statistical Association* 44 (247): 335–341.

Mikolov, Tomas, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. *Efficient Estimation of Word Representations in Vector Space.* Technical report. arXiv:1301.3781.

Minsky, M. L., and S. A. Papert. 1969. *Perceptrons.* Expanded edition 1990. MIT Press.

Mirza, M., and S. Osindero. 2014. *Conditional Generative Adversarial Nets.* Technical report. arXiv:1411.1784.

Miskin, J. W., and D. J. C. MacKay. 2001. "Ensemble learning for blind source separation." In *Independent Component Analysis: Principles and Practice,* edited by S. J. Roberts and R. M. Everson. Cambridge University Press.

Møller, M. 1993. "Efficient Training of Feed-Forward Neural Networks." PhD diss., Aarhus University, Denmark.

Montúfar, G. F., R. Pascanu, K. Cho, and Y. Bengio. 2014. "On the number of linear regions of deep neural networks." In *Proceedings of the International Conference on Neural Information Processing Systems (NeurIPS).* ArXiv:1402.1869.

Mordvintsev, Alexander, Christopher Olah, and Mike Tyka. 2015. *Inceptionism: Going Deeper into Neural Networks.* Google AI blog.

Murphy, Kevin P. 2022. *Probabilistic Machine Learning: An introduction.* MIT Press. probml.ai.

Murphy, Kevin P. 2023. *Probabilistic Machine Learning: Advanced Topics.* MIT Press. http://probml.github.io/book2.

Nakkiran, Preetum, Gal Kaplun, Yamini Bansal, Tristan Yang, Boaz Barak, and Ilya Sutskever. 2019. *Deep Double Descent: Where Bigger Models and More Data Hurt.* Technical report. arXiv:1912.02292.

Neal, R. M. 1993. *Probabilistic inference using Markov chain Monte Carlo methods.* Technical report CRG-TR-93-1. Department of Computer Science, University of Toronto, Canada.

Neal, R. M. 1999. "Suppressing random walks in Markov chain Monte Carlo using ordered overrelaxation." In *Learning in Graphical Models,* edited by Michael I. Jordan, 205–228. MIT Press.

Neal, R. M., and G. E. Hinton. 1999. "A new view of the EM algorithm that justifies incremental and other variants." In *Learning in Graphical Models,* edited by M. I. Jordan, 355–368. MIT Press.

Nelder, J. A., and R. W. M. Wedderburn. 1972. "Generalized linear models." *Journal of the Royal Statistical Society, A* 135:370–384.

Nesterov, Y. 2004. *Introductory Lectures on Convex Optimization: A Basic Course.* Kluwer.

Nichol, Alex, and Prafulla Dhariwal. 2021. *Improved Denoising Diffusion Probabilistic Models.* Technical report. arXiv:2102.09672.

Nichol, Alex, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. 2021. *GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models.* Technical report. arXiv:2112.10741.

Nocedal, J., and S. J. Wright. 1999. *Numerical Optimization.* Springer.

Noh, Hyeonwoo, Seunghoon Hong, and Bohyung Han. 2015. *Learning Deconvolution Network for Semantic Segmentation.* Technical report. arXiv:1505.04366.

Nowlan, S. J., and G. E. Hinton. 1992. "Simplifying neural networks by soft weight sharing." *Neural Computation* 4 (4): 473–493.

Ogden, R. T. 1997. *Essential Wavelets for Statistical Applications and Data Analysis.* Birkhäuser.

Oord, Aaron van den, Nal Kalchbrenner, and Koray Kavukcuoglu. 2016. *Pixel Recurrent Neural Networks.* Technical report. arXiv:1601.06759.

Oord, Aaron van den, Nal Kalchbrenner, Oriol Vinyals, Lasse Espeholt, Alex Graves, and Koray Kavukcuoglu. 2016. *Conditional Image Generation with PixelCNN Decoders.* Technical report. arXiv:1606.05328.

Oord, Aaron van den, Yazhe Li, and Oriol Vinyals. 2018. *Representation Learning with Contrastive Predictive Coding.* Technical report. arXiv:1807.03748.

Oord, Aaron van den, Oriol Vinyals, and Koray Kavukcuoglu. 2017. *Neural Discrete Representation Learning.* Technical report. arXiv:1711.00937.

OpenAI. 2023. *GPT-4 Technical Report.* Technical report. arXiv:2303.08774.

Opper, M., and O. Winther. 2000. "Gaussian processes and SVM: mean field theory and leave-one-out." In *Advances in Large Margin Classifiers,* edited by A. J. Smola, P. L. Bartlett, B. Schölkopf, and D. Shuurmans, 311–326. MIT Press.

Papamakarios, G., T. Pavlakou, and Iain Murray. 2017. "Masked Autoregressive Flow for Density Estimation." In *Proceedings of the International Conference on Neural Information Processing Systems (NeurIPS),* vol. 30.

Papamakarios, George, Eric Nalisnick, Danilo Jimenez Rezende, Shakir Mohamed, and Balaji Lakshminarayanan. 2019. *Normalizing Flows for Probabilistic Modeling and Inference.* Technical report. arXiv:1912.02762.

Parisi, Giorgio. 1981. "Correlation functions and computer simulations." *Nuclear Physics B* 180:378–384.

Pearl, J. 1988. *Probabilistic Reasoning in Intelligent Systems.* Morgan Kaufmann.

Pearlmutter, B. A. 1994. "Fast exact multiplication by the Hessian." *Neural Computation* 6 (1): 147–160.

Pearlmutter, B. A., and L. C. Parra. 1997. "Maximum likelihood source separation: a context-sensitive generalization of ICA." In *Advances in Neural Information Processing Systems,* edited by M. C. Mozer, M. I. Jordan, and T. Petsche, 9:613–619. MIT Press.

Pearson, Karl. 1901. "On lines and planes of closest fit to systems of points in space." *The London, Edinburgh and Dublin Philosophical Magazine and Journal of Science, Sixth Series* 2:559–572.

Phuong, Mary, and Marcus Hutter. 2022. *Formal Algorithms for Transformers.* Technical report. arXiv:2207.09238.

Prince, Simon J.D. 2020. *Variational autoencoders.* Https://www.borealisai.com/research-blogs/tutorial-5-variational-auto-encoders.

Prince, Simon J.D. 2023. *Understanding Deep Learning.* MIT Press. http://udlbook.com.

Radford, A., L. Metz, and S. Chintala. 2015. *Unsupervised representation learning with deep convolutional generative adversarial networks.* Technical report. arXiv:1511.06434.

Radford, Alec, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, *et al.* 2021. *Learning Transferable Visual Models From Natural Language Supervision.* Technical report. arXiv:2103.00020.

Radford, Alec, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. *Language Models are Unsupervised Multitask Learners.* Technical report. OpenAI.

Rakhimov, Ruslan, Denis Volkhonskiy, Alexey Artemov, Denis Zorin, and Evgeny Burnaev. 2020. *Latent Video Transformer.* Technical report. arXiv:2006.10704.

Ramachandran, P., B. Zoph, and Q. V. Le. 2017. *Searching for Activation Functions.* Technical report. arXiv:1710.05941v2.

Rao, C. R., and S. K. Mitra. 1971. *Generalized Inverse of Matrices and Its Applications.* Wiley.

Redmon, Joseph, Santosh Kumar Divvala, Ross B. Girshick, and Ali Farhadi. 2015. *You Only Look Once: Unified, Real-Time Object Detection.* Technical report. arxiv:1506.02640.

Ren, Shaoqing, Kaiming He, Ross B. Girshick, and Jian Sun. 2015. *Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks.* Technical report. arxiv:1506.01497.

Rezende, Danilo J, Shakir Mohamed, and Daan Wierstra. 2014. "Stochastic backpropagation and approximate inference in deep generative models." In *Proceedings of the 31st International Conference on Machine Learning (ICML-14),* 1278–1286.

Ricotti, L. P., S. Ragazzini, and G. Martinelli. 1988. "Learning of word stress in a sub-optimal second order backpropagation neural network." In *Proceedings of the IEEE International Conference on Neural Networks,* 1:355–361. IEEE.

Robert, C. P., and G. Casella. 1999. *Monte Carlo Statistical Methods.* Springer.

Rombach, Robin, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. 2021. *High-Resolution Image Synthesis with Latent Diffusion Models.* Technical report. arXiv:2112.10752.

Ronneberger, Olaf, Philipp Fischer, and Thomas Brox. 2015. "U-Net: Convolutional Networks for Biomedical Image Segmentation." In *Medical Image Computing and Computer-Assisted Intervention – MICCAI,* edited by N. Navab, J. Hornegger, W. Wells, and A. Frangi. Springer.

Rosenblatt, F. 1962. *Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms.* Spartan.

Roweis, S. 1998. "EM algorithms for PCA and SPCA." In *Advances in Neural Information Processing Systems,* edited by M. I. Jordan, M. J. Kearns, and S. A. Solla, 10:626–632. MIT Press.

Roweis, S., and Z. Ghahramani. 1999. "A unifying review of linear Gaussian models." *Neural Computation* 11 (2): 305–345.

Rubin, D. B., and D. T. Thayer. 1982. "EM algorithms for ML factor analysis." *Psychometrika* 47 (1): 69–76.

Rumelhart, D. E., G. E. Hinton, and R. J. Williams. 1986. "Learning internal representations by error propagation." In *Parallel Distributed Processing: Explorations in the Microstructure of Cognition,* edited by D. E. Rumelhart, J. L. McClelland, and the PDP Research Group, vol. 1: Foundations, 318–362. Reprinted in Anderson and Rosenfeld (1988). MIT Press.

Ruthotto, L., and E. Haber. 2021. *An introduction to deep generative modeling.* Technical report. arXiv:2103.05180.

Sagan, H. 1969. *Introduction to the Calculus of Variations.* Dover.

Saharia, Chitwan, William Chan, Huiwen Chang, Chris A. Lee, Jonathan Ho, Tim Salimans, David J. Fleet, and Mohammad Norouzi. 2021. *Palette: Image-to-Image Diffusion Models.* Technical report. arXiv:2111.05826.

Saharia, Chitwan, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily Denton, Seyed Kamyar Seyed Ghasemipour, *et al.* 2022. *Photorealistic Text-to-Image Diffusion Models*

with Deep Language Understanding. Technical report. arXiv:2205.11487.

Saharia, Chitwan, Jonathan Ho, William Chan, Tim Salimans, David J. Fleet, and Mohammad Norouzi. 2021. *Image Super-Resolution via Iterative Refinement.* Technical report. arXiv:2104.07636.

Santurkar, S., D. Tsipras, A. Ilyas, and A. Madry. 2018. *How does batch normalization help optimization?* Technical report. arXiv:1805.11604.

Satorras, Victor Garcia, Emiel Hoogeboom, and Max Welling. 2021. *E(n) Equivariant Graph Neural Networks.* Technical report. arXiv:2102.09844.

Schölkopf, B., and A. J. Smola. 2002. *Learning with Kernels.* MIT Press.

Schuhmann, Christoph, Richard Vencu, Romain Beaumont, Robert Kaczmarczyk, Clayton Mullis, Aarush Katta, Theo Coombes, Jenia Jitsev, and Aran Komatsuzaki. 2021. *LAION-400M: Open Dataset of CLIP-Filtered 400 Million Image-Text Pairs.* Technical report. arXiv:2111.02114.

Schuster, Mike, and Kaisuke Nakajima. 2012. "Japanese and Korean voice search." In *2012 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP),* 5149–5152.

Selvaraju, Ramprasaath R., Abhishek Das, Ramakrishna Vedantam, Michael Cogswell, Devi Parikh, and Dhruv Batra. 2016. *Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization.* Technical report. arXiv:1610.02391.

Sennrich, Rico, Barry Haddow, and Alexandra Birch. 2015. *Neural Machine Translation of Rare Words with Subword Units.* Technical report. arXiv:1508.07909.

Sermanet, Pierre, David Eigen, Xiang Zhang, Michael Mathieu, Rob Fergus, and Yann LeCun. 2013. *OverFeat: Integrated Recognition, Localization and Detection using Convolutional Networks.* Technical report. arXiv:1312.6229.

Shachter, R. D., and M. Peot. 1990. "Simulation Approaches to General Probabilistic Inference on Belief Networks." In *Uncertainty in Artificial Intelligence,* edited by P. P. Bonissone, M. Henrion, L. N. Kanal, and J. F. Lemmer, vol. 5. Elsevier.

Shannon, C. E. 1948. "A mathematical theory of communication." *The Bell System Technical Journal* 27 (3): 379–423 and 623–656.

Shen, Sheng, Zhen Dong, Jiayu Ye, Linjian Ma, Zhewei Yao, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. 2019. *Q-BERT: Hessian Based Ultra Low Precision Quantization of BERT.* Technical report. arXiv:1909.05840.

Simard, P., B. Victorri, Y. LeCun, and J. Denker. 1992. "Tangent prop – a formalism for specifying selected invariances in an adaptive network." In *Advances in Neural Information Processing Systems,* edited by J. E. Moody, S. J. Hanson, and R. P. Lippmann, 4:895–903. Morgan Kaufmann.

Simard, P. Y., D. Steinkraus, and J. Platt. 2003. "Best practice for convolutional neural networks applied to visual document analysis." In *Proceedings International Conference on Document Analysis and Recognition (ICDAR),* 958–962. IEEE Computer Society.

Simonyan, Karen, Andrea Vedaldi, and Andrew Zisserman. 2013. "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps." In *Computer Vision and Pattern Recognition.* ArXiv:1312.6034.

Simonyan, Karen, and Andrew Zisserman. 2014. *Very Deep Convolutional Networks for Large-Scale Image Recognition.* Technical report. arXiv:1409.1556.

Sirovich, L. 1987. "Turbulence and the Dynamics of Coherent Structures." *Quarterly Applied Mathematics* 45 (3): 561–590.

Sohl-Dickstein, Jascha, Eric A. Weiss, Niru Maheswaranathan, and Surya Ganguli. 2015. *Deep Unsupervised Learning using Nonequi-*

librium Thermodynamics. Technical report. arXiv:1503.03585.

Sønderby, C., J. Caballero, L. Theis, W. Shi, and F. Huszár. 2016. *Amortised MAP inference for image super-resolution.* Technical report. arXiv:1610.04490.

Song, Jiaming, Chenlin Meng, and Stefano Ermon. 2020. *Denoising Diffusion Implicit Models.* Technical report. arXiv:2010.02502.

Song, Yang, and Stefano Ermon. 2019. "Generative Modeling by Estimating Gradients of the Data Distribution." In *Advances in Neural Information Processing Systems,* 11895–11907. ArXiv:1907.05600.

Song, Yang, Sahaj Garg, Jiaxin Shi, and Stefano Ermon. 2019. "Sliced score matching: A scalable approach to density and score estimation." In *Uncertainty in Artificial Intelligence,* 204. ArXiv:1905.07088.

Song, Yang, and Diederik P. Kingma. 2021. *How to Train Your Energy-Based Models.* Technical report. arXiv:2101.03288.

Song, Yang, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. 2020. *Score-Based Generative Modeling through Stochastic Differential Equations.* Technical report. arXiv:2011.13456.

Srivastava, N., G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. 2014. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting." *Journal of Machine Learning Research* 15:1929–1958.

Stone, J. V. 2004. *Independent Component Analysis: A Tutorial Introduction.* MIT Press.

Sutskever, I., J. Martens, G. Dahl, and G. E. Hinton. 2013. "On the importance of initialization and momentum in deep learning." In *Proceedings of the International Conference on Machine Learning (ICML).*

Sutton, R. 2019. *The Bitter Lesson.* URL: incompleteideas.net/IncIdeas/BitterLesson.html.

Szegedy, Christian, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. 2013. *Intriguing properties of neural networks.* Technical report. arXiv:1312.6199.

Szeliski, R. 2022. *Computer Vision: Algorithms and Applications.* Second. Springer.

Tarassenko, L. 1995. "Novelty detection for the identification of masses in mamograms." In *Proceedings of the Fourth IEE International Conference on Artificial Neural Networks,* 4:442–447. IEE.

Tay, Yi, Mostafa Dehghani, Dara Bahri, and Donald Metzler. 2020. *Efficient Transformers: A Survey.* Technical report. arXiv:2009.06732.

Tibshirani, R. 1996. "Regression shrinkage and selection via the lasso." *Journal of the Royal Statistical Society, B* 58:267–288.

Tipping, M. E., and Christopher M. Bishop. 1997. *Probabilistic Principal Component Analysis.* Technical report NCRG/97/010. Neural Computing Research Group, Aston University.

Tipping, M. E., and Christopher M. Bishop. 1999. "Probabilistic Principal Component Analysis." *Journal of the Royal Statistical Society, Series B* 21 (3): 611–622.

Vapnik, V. N. 1995. *The nature of statistical learning theory.* Springer.

Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. *Attention Is All You Need.* Technical report. arXiv:1706.03762.

Veličković, Petar. 2023. *Everything is Connected: Graph Neural Networks.* Technical report. arXiv:2301.08210.

Veličković, Petar, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. 2017. *Graph Attention Networks.* Technical report. arXiv:1710.10903.

Vidakovic, B. 1999. *Statistical Modelling by Wavelets.* Wiley.

Vig, Jesse, Ali Madani, Lav R. Varshney, Caiming Xiong, Richard Socher, and Nazneen Fatema Rajani. 2020. *BERTology Meets Biology: Interpreting Attention in Protein Language Models.* Technical report. arXiv:2006.15222.

Vincent, P. 2011. "A connection between score matching and denoising autoencoders." *Neural Computation* 23:1661–1674.

Vincent, Pascal, Hugo Larochelle, Yoshua Bengio, and Pierre-Antoine Manzagol. 2008. "Extracting and Composing Robust Features with Denoising Autoencoders." In *Proceedings of the International Conference on Machine Learning (ICML).*

Walker, A. M. 1969. "On the asymptotic behaviour of posterior distributions." *Journal of the Royal Statistical Society, B* 31 (1): 80–88.

Wang, Chengyi, Sanyuan Chen, Yu Wu, Ziqiang Zhang, Long Zhou, Shujie Liu, Zhuo Chen, *et al.* 2023. *Neural Codec Language Models are Zero-Shot Text to Speech Synthesizers.* Technical report. arXiv:2301.02111.

Weisstein, E. W. 1999. *CRC Concise Encyclopedia of Mathematics.* Chapman / Hall, / CRC.

Welling, Max, and Yee Whye Teh. 2011. "Bayesian Learning via Stochastic Gradient Langevin Dynamics." In *Proceedings of the International Conference on Machine Learning (ICML).*

Williams, P. M. 1996. "Using neural networks to model conditional multivariate densities." *Neural Computation* 8 (4): 843–854.

Williams, R J. 1992. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine Learning* 8:229–256.

Winn, J., C. M. Bishop, T. Diethe, J. Guiver, and Y. Zaykov. 2023. *Model-Based Machine Learning*. Www.mbmlbook.com. Chapman / Hall.

Wolpert, D. H. 1996. "The lack of a-priori distinctions between learning algorithms." *Neural Computation* 8:1341–1390.

Wu, Zhirong, Yuanjun Xiong, Stella Yu, and Dahua Lin. 2018. *Unsupervised Feature Learning via Non-Parametric Instance-level Discrimination.* Technical report. arXiv:1805.01978.

Wu, Zonghan, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and Philip S. Yu. 2019. *A Comprehensive Survey on Graph Neural Networks.* Technical report. arXiv:1901.00596.

Yan, Wilson, Yunzhi Zhang, Pieter Abbeel, and Aravind Srinivas. 2021. *VideoGPT: Video Generation using VQ-VAE and Transformers.* Technical report. arXiv:2104.10157.

Yang, Ruihan, Prakhar Srivastava, and Stephan Mandt. 2022. *Diffusion Probabilistic Modeling for Video Generation.* Technical report. arXiv:2203.09481.

Yilmaz, Fatih Furkan, and Reinhard Heckel. 2022. *Regularization-wise double descent: Why it occurs and how to eliminate it.* Technical report. arXiv:2206.01378.

Yosinski, Jason, Jeff Clune, Anh Mai Nguyen, Thomas J. Fuchs, and Hod Lipson. 2015. *Understanding Neural Networks Through Deep Visualization.* Technical report. arXiv:1506.06579.

Yu, Jiahui, Xin Li, Jing Yu Koh, Han Zhang, Ruoming Pang, James Qin, Alexander Ku, Yuanzhong Xu, Jason Baldridge, and Yonghui Wu. 2021. *Vector-quantized Image Modeling with Improved VQGAN.* Technical report. arXiv:2110.04627.

Yu, Jiahui, Yuanzhong Xu, Jing Yu Koh, Thang Luong, Gunjan Baid, Zirui Wang, Vijay Vasudevan, *et al.* 2022. *Scaling Autoregressive Models for Content-Rich Text-to-Image Generation.* Technical report. arXiv:2206.10789.

Yu, Lili, Bowen Shi, Ramakanth Pasunuru, Benjamin Muller, Olga Golovneva, Tianlu Wang, Arun Babu, *et al.* 2023. *Scaling Autoregressive Multi-Modal Models: Pretraining and Instruction Tuning.* Technical report. arXiv:2309.02591.

Zaheer, Manzil, Satwik Kottur, Siamak Ravanbakhsh, Barnabas Poczos, Ruslan Salakhutdinov, and Alexander Smola. 2017. *Deep Sets.* Technical report. arXiv:1703.06114.

Zarchan, P., and H. Musoff. 2005. *Fundamentals of Kalman Filtering: A Practical Approach.* Second. AIAA.

Zeiler, Matthew D., and Rob Fergus. 2013. *Visualizing and Understanding Convolutional Networks.* Technical report. arXiv:1311.2901.

Zhang, Chiyuan, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. 2016. *Understanding deep learning requires rethinking generalization.* Technical report. arXiv:1611.03530.

Zhao, Wayne Xin, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, *et al.* 2023. *A Survey of Large Language Models.* Technical report. arXiv:2303.18223.

Zhou, Jie, Ganqu Cui, Shengding Hu, Zhengyan Zhang, Cheng Yang, Zhiyuan Liu, Lifeng Wang, Changcheng Li, and Maosong Sun. 2018. *Graph Neural Networks: A Review of Methods and Applications.* Technical report. arXiv:1812.08434.

Zhou, Y., and R. Chellappa. 1988. "Computation of optic flow using a neural network." In *International Conference on Neural Networks,* 71–78. IEEE.

Zhu, J-Y, T. Park, P. Isola, and A. Efros. 2017. *Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks.* Technical report. arXiv:1703.10593.