Index

Page numbers in **bold** indicate the primary source of information for the corresponding topic.

1×1 convolution, 296
1-of-K coding, 68, **135**, 460

acceptance criterion, **441**, 445, 448
activation, 17
activation function, 17, 158, 180, **182**
active constraint, 623
AdaGrad, 223
Adam optimization, 224
adaptive rejection sampling, 435
adjacency matrix, 410
adjoint sensitivity method, 556
adversarial attack, 306
aggregation, 415
aleatoric uncertainty, 23
AlexNet, 300
alpha family, 61
amortized inference, 572
ancestral sampling, 450
anchor, 191
annealed Langevin dynamics, 598
AR model, *see* autoregressive model
area under the ROC curve, 149
artificial intelligence, 1
attention, 358
attention head, 366
audio data, 399
auto-associative neural network, *see* autoencoder
autoencoder, 188, **563**
automatic differentiation, 22, 233, **244**
autoregressive flow, 552
autoregressive model, 5, **350**, 379
average pooling, 297

backpropagation, 19, **233**
backpropagation through time, 381
bag of words, 378
bagging, 278
base distribution, 547
basis function, 112, 158, 172, **172**
batch gradient descent, 214
batch learning, 117
batch normalization, 227
Bayes net, 326
Bayes' theorem, 28
Bayesian network, 326
Bayesian probability, 54
beam search, 386
Bernoulli distribution, **66**, 94
Bernoulli mixture model, 481
BERT, 388
bi-gram model, 379
bias, **39**, 125
bias parameter, **112**, 132, 180
bias–variance trade-off, 123
BigGAN, 539

© The Author(s), under exclusive license to Springer Nature Switzerland AG 2024
C. M. Bishop, H. Bishop, *Deep Learning*, https://doi.org/10.1007/978-3-031-45468-4

bijective function, 548
binomial distribution, 67
bits, 46
blind source separation, 514
blocked path, 339, **343**
boosting, 279
bootstrap, 278
bottleneck, 382
bounding box, 309
Box–Muller method, 432
byte pair encoding, 377

canonical correlation analysis, 501
canonical link function, 164
Cauchy distribution, 432
causal attention, 384
causality, 347
central differences, 239
central limit theorem, 71
ChatGPT, 394
child node, 249, **327**
Cholesky decomposition, 433
circular normal distribution, 89
classical probability, 54
classification, 3
CLIP, 192
co-parents, 348
codebook vector, 398, **465**
collider node, 341
combining models, 146
committee, 277
complete data set, 476
completing the square, 77
computer vision, 288
concave function, 52
concentration parameter, 92
condition number, 220
conditional entropy, 53
conditional expectation, 35
conditional independence, 146, **337**
conditional mixture model, 199
conditional probability, 27
conditional VAE, 576
conditioner, 552
confusion matrix, 147
continuous bag of words, 375
continuous normalizing flow, 557
contrastive divergence, 455
contrastive learning, 191
convex function, 51
convolution, **290**, 322
convolutional network, 287
correlation matrix, 503
cost function, 140
coupling flow, 549
coupling function, 552
covariance, 35
Cox's axioms, 54
cross attention, 390
cross-correlation, **292**, 322
cross-entropy error function, **160**, 162, 196
cross-validation, 14
cumulative distribution function, 32
curse of dimensionality, 172
curve fitting, 6
CycleGAN, 539

d-separation, 338, **343**, 479
DAG, *see* directed acyclic graph
data augmentation, 192, **257**
data compression, 465
DDIM, 594
DDPM, 581
decision, 120
decision boundary, 131, **139**
decision region, 131, **139**
decision surface, *see* decision boundary
decision theory, 120, **138**
decoder, 563
deep double descent, 268
deep learning, 20
deep neural networks, 20
deep sets, 417
DeepDream, 308
degrees of freedom, 495
denoising, 581
denoising autoencoder, 567
denoising diffusion implicit model, 594
denoising diffusion probabilistic model, 581
denoising score matching, 597

density estimation, 37, **65**
dequantization, 526
descendant node, 341
design matrix, 116
development set, 14
diagonal covariance matrix, 75
differential entropy, 50
diffusion kernel, 583
diffusion model, 581
Dirac delta function, 34
directed acyclic graph, 329
directed cycle, 329
directed factorization, 349
directed graph, 326
directed graphical model, 326
discriminant function, **132**, 143
discriminative model, **144**, 157, 346
disentangled representations, 542
distributed representation, 187
dot-product attention, 363
double descent, 268
dropout, 279

E step, **472**, 476
early stopping, 266
earth mover's distance, 538
ECM, *see* expectation conditional maximization
edge, **326**, 410
edge detection, 292
ELBO, *see* evidence lower bound
EM, *see* expectation maximization
embedding space, 188
embedding vector, 409
encoder, 563
energy function, 452
energy-based models, 452
ensemble methods, 277
entropy, 46
epistemic uncertainty, 23
epoch, 215
equality constraint, 623
equivariance, **259**, 292, 296, 371, 412
erf function, 164
error backpropagation, *see* backpropagation
error function, **8**, 55, 194, 210
Euler–Lagrange equations, 619
evaluation trace, 247
evidence lower bound, **485**, 516, 570, 588
expectation, 34
expectation conditional maximization, 489
expectation maximization, 470, **474**, 517, 519
expectation step, *see* E step
expectations, 430
explaining away, 343
exploding gradient, **227**, 382
exponential distribution, **34**, 431
exponential family, **94**, 156, 329
expression swell, 245

factor analysis, 513
factor graph, 327
factor loading, 513
false negative, 25
false positive, 25
fast gradient sign method, 306
fast R-CNN, 314
feature extraction, **20**, 113
feature map, 291
features, 179
feed-forward network, 172, **193**
feed-forward networks, 19
few-shot learning, **191**, 394
filter, 291
fine-tuning, 3, 22, **189**, 392
flow matching, 558
forward kinematics, 199
forward problem, 198
forward propagation, 235
foundation model, **22**, 358, 392, 409
frequentist probability, 54
fuel system, 341
fully connected graphical model, 328
fully convolutional network, 318
functional, 617

Gabor filters, 302
gamma distribution, 434
GAN, *see* generative adversarial network
gated recurrent unit, 382
Gaussian, 36, **70**
Gaussian mixture, 86, 200, 271, **466**

GEM, *see* generalized EM algorithm
generalization, 6
generalized EM algorithm, 489
generalized linear model, **158**, 165
generative adversarial network, 533
generative AI, 4
generative model, 4, **144**, 346, 533
generative pre-trained transformer, 6, **383**
geometric deep learning, 424
Gibbs sampling, 446
global minimum, 211
GNN, *see* graph neural network
GPT, *see* generative pre-trained transformer
GPU, *see* graphics processing unit
gradient descent, 209
graph attention network, 421
graph convolutional network, 414
graph neural network, 407
graph representation learning, 409
graphical model, 326
graphical model factorization, 329
graphics processing unit, **20**, 358
group theory, 256
guidance, 600

Hadamard product, 550
Hamiltonian Monte Carlo, 451
handwritten digit, 501
He initialization, 216
head-to-head path, 341
head-to-tail path, 340
Heaviside step function, 161
Hessian matrix, **211**, 242
Hessian outer product approximation, 243
heteroscedastic, 200
hidden Markov model, 380, **480**
hidden unit, **19**, 180
hidden variable, *see* latent variable
hierarchical representation, 187
histogram density estimation, 98
history of machine learning, 16
hold-out set, 14
homogeneous Markov chain, 443
Hooke's law, 520
Hutchinson's trace estimator, 557
hybrid Monte Carlo, 451
hyperparameter, 14

IAF, *see* inverse autoregressive flow
ICA, *see* independent component analysis
identifiability, 470
IID, *see* independent and identically distributed
image segmentation, 315
ImageNet data set, 299
importance sampling, **437**, 450
importance weight, 437
improper distribution, 33
improper prior, 263
inactive constraint, 623
incomplete data set, 476
independent and identically distributed, **37**, 344
independent component analysis, 514
independent factor analysis, 515
independent variables, 31
inductive bias, 19, **254**
inductive learning, **409**, 420
inequality constraint, 623
inference, 120, 138, **143**, 336
InfoNCE, 191
information theory, 46
instance discrimination, 192
internal covariate shift, 229
internal representation, 308
intersection-over-union, 310
intrinsic dimensionality, 496
invariance, 256, **256**, 297, 412
inverse autoregressive flow, 553
inverse kinematics, 199
inverse problem, 123, **198**, 254, 346
Iris data, 173
IRLS, *see* iterative reweighted least squares
isotropic covariance matrix, 75
iterative reweighted least squares, 160

Jacobian matrix, 44, **240**
Jensen's inequality, 52
Jensen–Shannon divergence, 544

K nearest neighbours, 103
K-means clustering algorithm, **460**, 480
Kalman filter, 353, **515**

Karush–Kuhn–Tucker conditions, 624
kernel density estimator, **100**, 596
kernel function, 101
kernel image, 291
KKT, *see* Karush-Kuhn-Tucker conditions
KL divergence, *see* Kullback–Leibler divergence
Kosambi–Karhunen–Loève transform, 497
Kullback–Leibler divergence, **51**, 486

Lagrange multiplier, 621
Lagrangian, 622
Langevin dynamics, 454
Langevin sampling, 455
language model, 382
Laplace distribution, 34
large language model, 5, 382, **390**
lasso, 264
latent class analysis, 481
latent diffusion model, 601
latent variable, 76, 335, **459**, 495
layer normalization, **229**, 369
LDM, *see* latent diffusion model
LDS, *see* linear dynamical system
leaky ReLU, 185
learning curve, 223, **266**
learning rate parameter, 214
learning to learn, 190
least-mean-squares algorithm, 118
least-squares GAN, 537
leave-one-out, 15
LeNet convolutional network, 299
Levenberg–Marquardt approximation, 244
likelihood function, **38**, 468
likelihood weighted sampling, 451
linear discriminant, 132
linear dynamical system, 515
linear independence, 610
linear regression, 6, **112**
linear-Gaussian model, 79, 332, **332**
linearly separable, 132
link, *see* edge
link function, **158**, 165
LLM, *see* large language model
LMS, *see* least-mean-squares algorithm
local minimum, 211
log odds, 151
logic sampling, 450
logistic regression, 159
logistic sigmoid, **95**, 113, 151, 159
logit function, 151
long short-term memory, 382
LoRA, *see* low-rank adaptation
loss function, **120**, 140
loss matrix, 142
lossless data compression, 465
lossy data compression, 465
low-rank adaptation, 392
LSGAN, *see* least-squares GAN
LSTM, *see* long short-term memory

M step, **472**, 477
macrostate, 48
MAE, *see* masked autoencoder
MAF, *see* masked autoregressive flow
Mahalanobis distance, 71
manifold, **177**, 522
MAP, *see* maximum a posteriori
marginal probability, 27
Markov blanket, **347**, 449
Markov boundary, *see* Markov blanket
Markov chain, **351**, 442
Markov chain Monte Carlo, 440
Markov model, 351
Markov random field, 327
masked attention, 384
masked autoencoder, 567
masked autoregressive flow, 553
max-pooling, 297
max-unpooling, 317
maximization step, *see* M step
maximum a posteriori, **56**, 477
maximum likelihood, **38**, 84, 115, 153
MCMC, *see* Markov chain Monte Carlo
MDN, *see* mixture density network
mean, 36
mean value theorem, 49
measure theory, 33
mel spectrogram, 399
message-passing, 414
message-passing neural network, 415

meta-learning, 190
Metropolis algorithm, 441
Metropolis–Hastings algorithm, 445
microstate, 48
mini-batches, 216
minimum risk, 145
Minkowski loss, 122
missing at random, **477**, 519
missing data, 519
mixing coefficient, 87
mixture component, 87
mixture density network, 198
mixture distribution, 459
mixture model, 459
mixture of Gaussians, 86, 200, 271, **466**
MLP, *see* multilayer perceptron
MNIST data, 495
mode collapse, 536
model averaging, 277
model comparison, 9
model selection, 14
moment, 37
momentum, 220
Monte Carlo dropout, 280
Monte Carlo sampling, 429
Moore-Penrose pseudo-inverse, *see* pseudo-inverse
MRF, *see* Markov random field
multi-class logistic regression, 161
multi-head attention, 366
multilayer perceptron, 18, **172**
multimodal transformer, 394
multimodality, 199
multinomial distribution, **70**, 95
multiplicity, 48
multitask learning, 190
mutual information, 54

n-gram model, 379
naive Bayes model, 147, **344**, 378
nats, 47
natural language processing, 374
natural parameter, 94
nearest-neighbours, 103
neocognitron, 302
Nesterov momentum, 221
neural ordinary differential equation, 554
neuroscience, 302
NLP, *see* natural language processing
no free lunch theorem, 255
node, **326**, 410
noise, 23
noiseless coding theorem, 47
noisy-OR, 354
non-identifiability, 513
non-max suppression, 314
nonparametric methods, 66, **98**
normal distribution, *see* Gaussian
normal equations, 116
normalized exponential, *see* softmax function
novelty detection, 144

object detection, 308
observed variable, 335
Old Faithful data, 86
on-hot encoding, *see* 1-of-K encoding
one-shot learning, 191
one-versus-one classifier, 134
one-versus-the-rest classifier, 134
online gradient descent, 215
online learning, 117
ordered over-relaxation, 449
outer product approximation, 244
outlier, **137**, 144, 164
over-fitting, **10**, 123, 470
over-relaxation, 449
over-smoothing, 422

padding, 294
parameter sharing, **270**, 331
parameter shrinkage, 118
parameter tying, *see* parameter sharing
parent node, 247, **327**
partition function, 452
Parzen estimator, *see* kernel density estimator
Parzen window, 101
PCA, *see* principal component analysis
perceptron, 17
periodic variables, 89
permutation matrix, 411
PixelCNN, 397
PixelRNN, 397

plate, 334
polynomial curve fitting, 6
pooling, 296
positional encoding, 371
positive definite covariance, 72
positive definite matrix, 615
posterior collapse, 577
posterior probability, 31
power method, 498
pre-activation, 17
pre-processing, 20
pre-training, **189**, 392
precision matrix, 77
precision parameter, 36
predictive distribution, **42**, 120
prefix prompt, 394
principal component analysis, **497**, 506, 565
principal subspace, 497
prior, 263
prior knowledge, 19, **255**
prior probability, **31**, 145
probabilistic graphical model, *see* graphical model
probabilistic PCA, 506
probability, 25
probability density, 32
probability theory, 23
probit function, 164
probit regression, 163
product rule of probability, 26, **28**, 326
prompt, **394**, 601
prompt engineering, 394
proposal distribution, **433**, 437, 441
pseudo-inverse, **116**, 136
pseudo-random numbers, 430

quadratic discriminant, 153

radial basis functions, 179
random variable, 26
raster scan, 397
readout layer, 419
real NVP normalizing flow, 549
receiver operating characteristic, *see* ROC curve
receptive field, **290**, 416
recurrent neural network, 380
regression, 3
regression function, 121
regularization, 12, **253**
regularized least squares, 118
reject option, **142**, 145
rejection sampling, 433
relative entropy, 51
reparameterization trick, 574
representation learning, 22, **188**
residual block, 275
residual connection, 22, **274**
residual network, 275
resnet, *see* residual network
responsibility, 88, **468**
RLHF, 394
RMS error, *see* root-mean-square error
RMSProp, 223
RNN, *see* recurrent neural network
robot arm, 198
robustness, 137
ROC curve, 148
root-mean-square error, 10

saliency map, 305
same convolution, 294
sample mean, 39
sample variance, 39
sampling, 429
sampling-importance-resampling, 439
scale invariance, 256
scaled self-attention, 366
scaling hypothesis, 358
Schur complement, 79
score function, 455, 594
score matching, 594
self-attention, 362
self-supervised learning, **5**, 375
semi-supervised learning, 420
sequential estimation, 85
sequential gradient descent, 118
sequential learning, 117
SGD, *see* stochastic gradient descent
shared parameters, *see* parameter sharing
shared weights, 292
shattered gradients, 274
shrinkage, 13

sigmoid, *see* logistic sigmoid
singular value decomposition, 117
SIR, *see* sampling-importance-resampling
skip-grams, 375
skip-layer connections, 274
sliding window, 311
smoothing parameter, 100
soft ReLU, 185
soft weight sharing, 271
softmax function, 96, **152**, 197, 201, 363
softplus activation function, 185
sparse autoencoders, 566
sparse connections, 292
sparsity, 264
sphering, 504
standard deviation, 36
standardizing, 462, **503**
state-space model, 352
statistical bias, *see* bias
statistical independence, *see* independent variables
steepest descent, 214
Stein score, *see* score function
Stirling's approximation, 48
stochastic, 8
stochastic differential equation, 598
stochastic gradient descent, 19, **214**, 215
stochastic variable, 26
strided convolution, 294
strides, 311
structured data, 287, **407**
style transfer, 320
sufficient statistics, 67, 69, 84, **97**
sum rule of probability, 26, **28**, 326
sum-of-squares error, **8**, 41, 136
supervised learning, **3**, 420
support vector machine, 179
SVD, *see* singular value decomposition
SVM, *see* support vector machine
swish activation function, 205
symmetry, 256
symmetry breaking, 216

tail-to-tail path, 339
tangent propagation, 258
temperature, 387
tensor, **194**, 295
test set, 10, **14**
text-to-speech, 400
tied parameters, *see* parameter sharing
token, 360
tokenization, 377
training set, 3
transductive, 409, **419**
transductive learning, 420
transfer learning, 3, **189**, 218, 388
transformers, 357
transition probability, 443
translation invariance, 256
transpose convolution, 318
tri-gram model, 379
TTS, *see* text-to-speech

U-net, 319
undetermined multiplier, *see* Lagrange multiplier
undirected graphical model, 327
uniquenesses, 513
universal approximation theorems, 182
unobserved variable, *see* latent variable
unsupervised learning, **4**, 188
utility function, 140

VAE, *see* variational autoencoder
valid convolution, 294
validation set, 14
vanishing gradient, **227**, 382
variance, **35**, 36, 125
variational autoencoder, 569
variational inference, 485
variational lower bound, *see* evidence lower bound
vector quantization, 398, **465**
vertex, *see* node
vision transformer, 395
von Mises distribution, 89
voxel, 289

Wasserstein distance, 538
Wasserstein GAN, 538
wavelets, 114
weakly supervised, 192
weight decay, 13, **260**
weight parameter, **17**, 180

weight sharing, *see* parameter sharing
weight vector, 132
weight-space symmetry, 185
WGAN, *see* Wasserstein GAN
whitening, 502
Woodbury identity, 610
word embedding, 375
word2vec, 375
wrapped distribution, 94

Yellowstone National Park, 86